

Installed Packages

```
In [ ]: import sys
!{sys.executable} -m pip install contractions
!{sys.executable} -m pip install gensim==4.2.0
!pip install scikit-learn
!pip install torch torchvision torchaudio

Looking in indexes: https://pypi.org/simple, (https://pypi.org/simple,) https://us-python.pkg.dev/colab-wheels/public/simple/ (https://us-python.pkg.dev/colab-wheels/public/simple/)
Collecting contractions
  Downloading contractions-0.1.73-py2.py3-none-any.whl (8.7 kB)
Collecting textsearch>=0.0.21
  Downloading textsearch-0.0.24-py2.py3-none-any.whl (7.6 kB)
Collecting anyascii
  Downloading anyascii-0.3.1-py3-none-any.whl (287 kB)
    _____ 287.5/287.5 KB 7.5 MB/s eta 0:00:00
Collecting pyahocorasick
  Downloading pyahocorasick-2.0.0-cp38-cp38-manylinux_2_5_x86_64.manylinux1_x86_64.whl (104 kB)
    _____ 104.5/104.5 KB 8.5 MB/s eta 0:00:00
Installing collected packages: pyahocorasick, anyascii, textsearch, contractions
Successfully installed anyascii-0.3.1 contractions-0.1.73 pyahocorasick-2.0.0 textsearch-0.0.24
Looking in indexes: https://pypi.org/simple, (https://pypi.org/simple,) https://us-python.pkg.dev/colab-wheels/public/simple/ (https://us-python.pkg.dev/colab-wheels/public/simple/)
Collecting gensim==4.2.0
  Downloading gensim-4.2.0-cp38-cp38-manylinux_2_12_x86_64.manylinux2010_x86_64.whl (24.1 MB)
    _____ 24.1/24.1 MB 28.9 MB/s eta 0:00:00
- - - - -
```

```
In [ ]: ## Importing and installing libraries

import numpy as np
import copy
import pandas as pd
import warnings
import re
import sys
import nltk
from gensim.models import Word2Vec
from nltk.corpus import stopwords
import string
from torch import nn
import torch.nn as nn
import torch.nn.functional as F
import torch
from torch.nn import CrossEntropyLoss, Softmax, Linear
from torch.optim import SGD, Adam
from sklearn.metrics.pairwise import cosine_similarity
from torch.optim.lr_scheduler import ReduceLROnPlateau
from nltk.stem import WordNetLemmatizer
from gensim.models import KeyedVectors
from gensim import utils
from scipy.sparse import hstack
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Perceptron
from sklearn.metrics import classification_report
from sklearn import svm
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import LinearSVC
from statistics import mean
from os import path
import os.path
import gensim
import gensim.downloader
from sklearn.svm import LinearSVC

nltk.download('punkt')

warnings.filterwarnings('ignore')

import contractions

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
```

1. Dataset Generation

```
In [ ]: from google.colab import drive
drive.mount('/content/drive/')
%cd /content/drive/My Drive/Colab Notebooks/
```

Mounted at /content/drive/
/content/drive/My Drive/Colab Notebooks

```
In [ ]: #fields required in the balanced dataframe from the original dataset
input_column=["review_body", "star_rating"]

#reading the original dataset to filter the columns that are required
input_df =pd.read_csv('./amazon_reviews_us_Beauty_v1_00.tsv',usecols=input_column,sep='\t',error_bad_lines=False)
```

```
In [ ]: #Creating 3 different classes to get 20000 data from each class to avoid computational burden

class_one_df =(input_df[(input_df['star_rating'] == 1) | (input_df['star_rating'] == 2) ]).sample(n=20000)
class_one_df['class']=1

class_two_df =(input_df[(input_df['star_rating'] == 3)]).sample(n=20000)
class_two_df['class']=2

class_three_df =(input_df[(input_df['star_rating'] == 4) | (input_df['star_rating'] == 5) ]).sample(n=20000)
class_three_df['class']=3

#Combining all the data received from each class into a single balanced dataframe

amazon_balanced_df = pd.concat([class_one_df, class_two_df, class_three_df])

#Resetting the index as we have retrieved different data according to the classes created.
#Therefore, we will have irregular or unsorted index keys.
#We will reset the index to the new and incremental values from 0

amazon_balanced_df = amazon_balanced_df.reset_index(drop=True)

# Created a new dataframe consisting of the two columns (star_rating and review_body)
#along with class one assigned to them on the basis of star_rating. We are also resetting the index
```

Data Cleaning

Handling null values

```
In [ ]: #We are changing all null values to an empty string

amazon_balanced_df = amazon_balanced_df.fillna('')
```

```
In [ ]: #Uncleaned data copy
amazon_df=amazon_balanced_df.copy()
```

Convert all reviews into lowercase

```
In [ ]: # Converting all review body into lowercase

amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.lower()
```

Remove the HTML from the reviews

```
In [ ]: # Removing all the html tags from each review body

amazon_balanced_df['review_body']=amazon_balanced_df['review_body'].apply(lambda x : re.sub('<.*?>','',str(x)))
```

Remove the URLs from the reviews

```
In [ ]: # Removing all the URLs from each review body

amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].apply(lambda y: re.split('https://\/.*', s
```

Remove non-alphabetical characters

```
In [ ]: # Removing all the non alphabetic chaarcters(symbols, numbers) from each review body

amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].apply(lambda z: " ".join([re.sub('[^A-Za-z',
```

Remove extra spaces

```
In [ ]: # Will remove leading and trailing spaces

amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.strip()
```

Perform contractions on the review_body

```
In [ ]: ## This will elongate the short form used in sentences like (I'll ---> I will)

amazon_balanced_df['without_contraction'] = amazon_balanced_df['review_body'].apply(lambda a: [contractions.fix
amazon_balanced_df['review_body'] = [' '.join(map(str, x)) for x in amazon_balanced_df['without_contraction']]
```

Remove Punctuations

```
In [ ]: amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.replace(r'[\w\s]+', '')
```

2. Word Embedding

(a) Downloading pretrained word2vec-google-news-300

```
In [ ]: word2vec_model = gensim.downloader.load('word2vec-google-news-300')

[=====] 99.7% 1657.5/1662.8MB downloaded
```

```
In [ ]: word2vec_model.save('gensim2.kv')
```

```
In [ ]: print(cosine_similarity([word2vec_model['queen']], [word2vec_model['king'] - word2vec_model['man'] + word2vec_model['king']]))
print(cosine_similarity([word2vec_model['queen']], [word2vec_model['king']]))

[[0.7300518]]
[[0.6510957]]
```

```
In [ ]: word2vec_model.most_similar("good")
```

```
Out[18]: [('great', 0.7291510105133057),
 ('bad', 0.7190051078796387),
 ('terrific', 0.6889115571975708),
 ('decent', 0.6837348341941833),
 ('nice', 0.6836092472076416),
 ('excellent', 0.644292950630188),
 ('fantastic', 0.6407778263092041),
 ('better', 0.6120728850364685),
 ('solid', 0.5806034803390503),
 ('lousy', 0.576420247554779)]
```

```
In [ ]: word2vec_model.similarity(w1="daughter", w2="sister")
```

```
Out[19]: 0.7814771
```

(b) Training word2vec model on our own dataset

```
In [ ]: class dataEmbed:
    def __init__(self, data_set):
        self.data_set = data_set

    def __iter__(self):
        for x in self.data_set:
            yield utils.simple_preprocess(x)
```

```
In [ ]: sentence_embed = dataEmbed(amazon_balanced_df.review_body)
# window=13
# vector_size=300
# min_count=9
embed_word2vec = Word2Vec(sentences=sentence_embed, vector_size=300, min_count=9, window=13)
model = embed_word2vec.wv
```

```
In [ ]: print(cosine_similarity([model['queen']], [model['king'] - model['man'] + model['woman']]))
print(cosine_similarity([model['queen']], [model['king']]))

[[0.1676019]]
[[0.40402395]]
```

```
In [ ]: model.most_similar("good")
```

```
Out[23]: [('great', 0.7492009997367859),
('decent', 0.6979432106018066),
('nice', 0.6569323539733887),
('fantastic', 0.600950300693512),
('ok', 0.5753679275512695),
('bad', 0.5533730387687683),
('okay', 0.5282017588615417),
('alright', 0.5122868418693542),
('awesome', 0.5109883546829224),
('high', 0.4818119406700134)]
```

```
In [ ]: model.similarity(w1="daughter", w2="sister")
```

```
Out[24]: 0.8866891
```

3. Simple Models

Split the data into 80% train and 20% test

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(amazon_balanced_df['review_body'], amazon_balanced_df['class'],
```

```
In [ ]: # print("Train Size ", X_train.shape)
# print("Test Size ", X_test.shape)
```

TF-IDF

```
In [ ]: Tf_Id_Vector = TfidfVectorizer(
ngram_range=(1, 2),
analyzer='word',
token_pattern=r'\w{1,}',
strip_accents='unicode',
max_features=10000,
stop_words='english'
)
```

```
In [ ]: X_tf_id_train = Tf_Id_Vector.fit_transform(X_train)
X_tf_id_test = Tf_Id_Vector.transform(X_test)
print("Train feature names", Tf_Id_Vector.get_feature_names_out())
```

```
Train feature names ['aa' 'ability' 'able' ... 'zipper' 'zits' 'zone']
```

Perceptron

```
In [ ]: perceptron = Perceptron(
penalty= 'l1', #Provided to model on wrong prediction
alpha=0.000005, #Strength of the L2 regularization term
max_iter=1500, #Maximum number of iterations over the data
tol=1e-1, #Tolerance for the optimization or the criterion for stopping
)
perceptron.fit(X_tf_id_train , y_train)
#prediction through X_test data

y_test_prediction=perceptron.predict(X_tf_id_test)
#Preparing the report by comparing actual value and predicted value
report=classification_report(y_test, y_test_prediction)
print("\n Values for Perceptron Model")
print(report)
```

Values for Perceptron Model					
	precision	recall	f1-score	support	
1	0.61	0.69	0.65	4029	
2	0.52	0.51	0.52	3990	
3	0.73	0.67	0.70	3981	
accuracy			0.62	12000	
macro avg	0.62	0.62	0.62	12000	
weighted avg	0.62	0.62	0.62	12000	

SVM

```
In [ ]: svm = LinearSVC(
C=0.10, #Regularization parameter. Default is 1.
penalty='l2', #Norm of Penalty
tol=1e-1, #tolerance of Stopping criteria. default is 1e-3
class_weight="balanced", #adjust and provides weight to each class
max_iter=1000, #Hard limit on iterations within solver, or -1 for no limit.
random_state=1, #Controls the pseudo random number generation for shuffling the data for probability estimates.
loss='squared_hinge', #Specifies the Loss Function
dual=False, #Selects the algorithm to either the dual or primal optimization
fit_intercept=False, #Whether to calculate intercept for this model
)
svm.fit(X_tf_id_train, y_train)
#prediction through X_test data
y_test_prediction_svm=svm.predict(X_tf_id_test)
#Preparing the report by comparing actual value and predicted value
report_svm=classification_report(y_test, y_test_prediction_svm)
print("\n Values for SVM Model")
print(report_svm)
```

Values for SVM Model					
	precision	recall	f1-score	support	
1	0.69	0.70	0.70	4029	
2	0.61	0.56	0.58	3990	
3	0.74	0.79	0.76	3981	
accuracy			0.68	12000	
macro avg	0.68	0.68	0.68	12000	
weighted avg	0.68	0.68	0.68	12000	

Process to extract word2vec embeddings

```
In [ ]: embedding_space = []
        for i in range(60000):
            vectorWord = np.zeros((1,300))
            listword = amazon_df['review_body'][i].split(" ")
            for word in listword:
                if word in word2vec_model.key_to_index:
                    np.reshape(word2vec_model[word], (1, 300))
                    vectorWord += word2vec_model[word]
                else:
                    vectorWord += np.zeros((1,300))
            avg_wordVec = vectorWord/len(listword)
            embedding_space.append(avg_wordVec)

embedding_dataset = np.array(embedding_space)
print(embedding_dataset.shape)
embedding_dataset = embedding_dataset.reshape(embedding_dataset.shape[0], embedding_dataset.shape[2])

(60000, 1, 300)
```

```
In [ ]: embedding_space_concat = []
        for i in range(60000):
            vectorWord = [] # change the size of the vector
            listword = amazon_df['review_body'][i].split(" ")
            for item in listword[:20]:
                if item in word2vec_model:
                    x=np.reshape(word2vec_model[item], (1, 300))
                    vectorWord.append(x)
            vectorWord=vectorWord[1:]
            if len(vectorWord) < 20:
                di = 20 - len(vector_word)
                vectorWord += [np.zeros((1, 300))] * di

            embedding_space_concat.append(vectorWord)
embedding_space_concat=np.array(embedding_space_concat)
embedding_dataset=embedding_space_concat.reshape(embedding_space_concat.shape[0], embedding_space_concat.shape[1], embedding_space_concat.shape[2])

(60000, 20, 300)
```

```
In [ ]: A_train, A_test, B_train, B_test = train_test_split(embedding_dataset, amazon_df['class'], test_size=0.20, random_state=42)

B_train = B_train.reset_index(drop=True)
B_test = B_test.reset_index(drop=True)

print(A_train.shape, A_test.shape, B_train.shape, B_test.shape)

(48000, 300) (12000, 300) (48000,) (12000,)
```

Perceptron

```
In [ ]: perceptron = Perceptron(
        penalty= 'l1', #Provided to model on wrong prediction
        alpha=0.000005, #Strength of the L2 regularization term
        max_iter=1500, #Maximum number of iterations over the data
        tol=1e-1, #Tolerance for the optimization or the criterion for stopping
    )
perceptron.fit(A_train , B_train)
#prediction through X_test data

B_test_prediction=perceptron.predict(A_test)
#Preparing the report by comparing actual value and predicted value
report=classification_report(B_test, B_test_prediction)
print("\n Values for Perceptron Model")
print(report)
```

Values for Perceptron Model				
	precision	recall	f1-score	support
1	0.67	0.50	0.58	4000
2	0.58	0.32	0.41	4000
3	0.51	0.87	0.65	4000
accuracy			0.56	12000
macro avg	0.59	0.56	0.54	12000
weighted avg	0.59	0.56	0.54	12000

SVM

```
In [ ]: svm = LinearSVC(
C=0.10, #Regularization parameter. Default is 1.
penalty='l2', #Norm of Penalty
tol=1e-1, #tolerance of Stopping criteria. default is 1e-3
class_weight="balanced", #adjust and provides weight to each class
max_iter=1000, #Hard limit on iterations within solver, or -1 for no limit.
random_state=1, #Controls the pseudo random number generation for shuffling the data for probability estimates.
loss='squared_hinge', #Specifies the Loss Function
dual=False, #Selects the algorithm to either the dual or primal optimization
fit_intercept=False, #Weather to calculate intercept for this model
)
svm.fit(A_train , B_train)
#prediction through X_test data
B_test_prediction_svm=svm.predict(A_test)
#Preparing the report by comparing actual value and predicted value
report_svm=classification_report(B_test, B_test_prediction_svm)
print("\n Values for SVM Model")
print(report_svm)
```

Values for SVM Model					
	precision	recall	f1-score	support	
1	0.60	0.66	0.63	4000	
2	0.56	0.49	0.53	4000	
3	0.67	0.70	0.68	4000	
accuracy			0.62	12000	
macro avg	0.61	0.62	0.61	12000	
weighted avg	0.61	0.62	0.61	12000	

4. Feedforward Neural Networks

```
In [ ]: from torch.utils.data import Dataset, DataLoader
```

```
In [ ]: #Creating a dataloader using torch
class dataloader(torch.utils.data.Dataset):
    def __init__(self, dataset_record, label_record):
        self.dataset = dataset_record
        self.labels = label_record

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, index):
        dataset = self.dataset[index]
        labels = self.labels[index]

        return dataset, labels
```

```
In [ ]: #Creating classes to define the architecure
class feedForward(nn.Module):
    def __init__(self, output_size, input_size):
        super(feedForward, self).__init__()
        self.layer1 = nn.Linear(input_size, 100)
        self.relu1 = nn.ReLU()
        self.layer2 = nn.Linear(100, 10)
        self.relu2 = nn.ReLU()
        self.layer3 = nn.Linear(10, output_size)

    def forward(self, x):
        return self.layer3(self.relu2(self.layer2(self.relu1(self.layer1(x)))))
```

```
In [ ]: fnn=feedForward(3,300)
fnn
```

```
Out[42]: feedForward(
  (layer1): Linear(in_features=300, out_features=100, bias=True)
  (relu1): ReLU()
  (layer2): Linear(in_features=100, out_features=10, bias=True)
  (relu2): ReLU()
  (layer3): Linear(in_features=10, out_features=3, bias=True)
)
```

(a) Testing split of Multi layer perceptron and fetching accuracy of FNN

```
In [ ]: # Convert A_train and A_test to float32
A_word2vec_train = A_train.astype(np.float32)
A_word2vec_test  = A_test.astype(np.float32)

# Subtract 1 from B_train and B_test values
B_train = B_train - 1
B_test  = B_test  - 1

# Create PyTorch DataLoader objects for the training and testing sets
train_dataset = dataloader(A_word2vec_train, B_train)
train_set     = torch.utils.data.DataLoader(train_dataset, batch_size=50)

test_dataset = dataloader(A_word2vec_test, B_test)
test_set     = torch.utils.data.DataLoader(test_dataset, batch_size=50)
```

```
In [ ]: from sklearn.metrics import accuracy_score, f1_score
```



```

In [ ]: def train(reviews_dataloader_train, reviews_dataloader_test, model, num_epochs, concat=False, rnn=False, gru=False):
    y_pred_label_train = []
    y_true_label_train = []
    y_pred_label_test = []
    y_true_label_test = []

    # Set the device for the model
    # device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    # model.to(device)

    # Define the loss function and optimizer
    criterion = nn.CrossEntropyLoss()
    optimizer = Adam(model.parameters(), lr=0.001)
    softmax = Softmax(dim=1)

    # Define the scheduler
    scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=5, gamma=0.1)

    # Keep track of the best model
    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0

    # Keep track of the previous loss
    loss_min = prev_loss

    # Train the model
    for epoch in range(num_epochs):
        print('\n Epoch: {}'.format(epoch))

        # print(reviews_dataloader_train)
        for j, (x, y) in enumerate(reviews_dataloader_train):
            y_pred = model(x)
            y_pred_label_train.append(torch.argmax(softmax(y_pred.detach()), axis=1))
            y_true_label_train.append(y.detach())
            loss = criterion(y_pred, y)

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            # if j % 100 == 0:
            #     print('Epoch {:03} Batch {:03}/{:03} Loss: {:.4f}'.format(epoch, j, len(reviews_dataloader_train), loss.item()))

        # Evaluate the model on the test set
        with torch.no_grad():
            for x, y in reviews_dataloader_test:
                y_pred = model(x)
                y_pred_label_test.append(torch.argmax(softmax(y_pred.detach()), axis=1))
                y_true_label_test.append(y.detach())

        # Calculate accuracy and f1-score
        y_pred_train = torch.cat(y_pred_label_train)
        y_true_train = torch.cat(y_true_label_train)
        y_pred_test = torch.cat(y_pred_label_test)
        y_true_test = torch.cat(y_true_label_test)

        train_acc = accuracy_score(y_true_train.cpu().numpy(), y_pred_train.cpu().numpy())
        test_acc = accuracy_score(y_true_test.cpu().numpy(), y_pred_test.cpu().numpy())
        train_f1 = f1_score(y_true_train.cpu().numpy(), y_pred_train.cpu().numpy(), average='macro')
        test_f1 = f1_score(y_true_test.cpu().numpy(), y_pred_test.cpu().numpy(), average='macro')

        print('Epoch: {:03}, Loss: {:.4f}, Train Acc: {:.4f}, Test Acc: {:.4f}'.format(epoch, loss.item(), train_acc, test_acc))

        # Update the learning rate
        scheduler.step()

        # Save the best model based on test accuracy
        if test_acc > best_acc:
            best_acc = test_acc
            best_model_wts = copy.deepcopy(model.state_dict())

        # Save the model checkpoint
        # if loss.item() < loss_min:
        #     print(f'Loss decreased from {loss_min:.4f} to {loss.item():.4f}. Saving model...')
        #     torch.save(model.state_dict(), 'model_checkpoint.pt')
        #     loss

```

```
In [ ]: train(train_set, test_set, fnn, 20)
```

```
Epoch: 0
Epoch: 000, Loss: 0.8441, Train Acc: 0.5594, Test Acc: 0.6198

Epoch: 1
Epoch: 001, Loss: 0.8092, Train Acc: 0.5902, Test Acc: 0.6267

Epoch: 2
Epoch: 002, Loss: 0.7992, Train Acc: 0.6043, Test Acc: 0.6308

Epoch: 3
Epoch: 003, Loss: 0.7932, Train Acc: 0.6129, Test Acc: 0.6334

Epoch: 4
Epoch: 004, Loss: 0.7883, Train Acc: 0.6187, Test Acc: 0.6351

Epoch: 5
Epoch: 005, Loss: 0.7794, Train Acc: 0.6241, Test Acc: 0.6365

Epoch: 6
Epoch: 006, Loss: 0.7783, Train Acc: 0.6281, Test Acc: 0.6376

Epoch: 7
Epoch: 007, Loss: 0.7764, Train Acc: 0.6312, Test Acc: 0.6384

Epoch: 8
Epoch: 008, Loss: 0.7749, Train Acc: 0.6337, Test Acc: 0.6391

Epoch: 9
Epoch: 009, Loss: 0.7736, Train Acc: 0.6357, Test Acc: 0.6396

Epoch: 10
Epoch: 010, Loss: 0.7743, Train Acc: 0.6374, Test Acc: 0.6399

Epoch: 11
Epoch: 011, Loss: 0.7742, Train Acc: 0.6388, Test Acc: 0.6402

Epoch: 12
Epoch: 012, Loss: 0.7740, Train Acc: 0.6400, Test Acc: 0.6404

Epoch: 13
Epoch: 013, Loss: 0.7739, Train Acc: 0.6410, Test Acc: 0.6407

Epoch: 14
Epoch: 014, Loss: 0.7737, Train Acc: 0.6419, Test Acc: 0.6408

Epoch: 15
Epoch: 015, Loss: 0.7739, Train Acc: 0.6427, Test Acc: 0.6410

Epoch: 16
Epoch: 016, Loss: 0.7740, Train Acc: 0.6434, Test Acc: 0.6412

Epoch: 17
Epoch: 017, Loss: 0.7741, Train Acc: 0.6440, Test Acc: 0.6414

Epoch: 18
Epoch: 018, Loss: 0.7742, Train Acc: 0.6445, Test Acc: 0.6415

Epoch: 19
Epoch: 019, Loss: 0.7742, Train Acc: 0.6450, Test Acc: 0.6417
```

```
In [ ]: ### To concatenate first 10 Word2Vec vectors for each review as the input feature
embedding_space_concat = []
for i in range(60000):
    vectorWord = np.zeros((1,300*10)) # change the size of the vector
    listword = amazon_df['review_body'][i].split(" ")
    for j, word in enumerate(listword):
        if j < 10: # only consider the first 10 words
            if word in word2vec_model.key_to_index:
                np.reshape(word2vec_model[word], (1, 300))
                vectorWord[0, j*300:(j+1)*300] = word2vec_model[word] # concatenate the vector
            else:
                vectorWord[0, j*300:(j+1)*300] = np.zeros((1,300))
    embedding_space_concat.append(vectorWord)

embedding_dataset_concat = np.array(embedding_space_concat)
print(embedding_dataset_concat.shape)
embedding_dataset_concat = embedding_dataset_concat.reshape(embedding_dataset_concat.shape[0], embedding_dataset_concat.shape[1]*10)
(60000, 1, 3000)
```

```
In [ ]: P_train, P_test, Q_train, Q_test = train_test_split(embedding_space_concat, amazon_df['class'], test_size=0.20,

Q_train = Q_train.reset_index(drop=True)
Q_test = Q_test.reset_index(drop=True)

P_train= np.array(P_train)
P_test= np.array(P_test)
```

```
In [ ]: type(P_train)
```

```
Out[68]: numpy.ndarray
```

```
In [ ]: # Convert A_train and A_test to float32
P_word2vec_train = P_train.astype(np.float32)
P_word2vec_test  = P_test.astype(np.float32)

# Subtract 1 from B_train and B_test values
Q_train = Q_train - 1
Q_test = Q_test - 1

# Create PyTorch DataLoader objects for the training and testing sets
train_datasetC = dataloader(P_word2vec_train, Q_train)
train_setC = torch.utils.data.DataLoader(train_datasetC, batch_size=50)

test_datasetC = dataloader(P_word2vec_test, Q_test)
test_setC = torch.utils.data.DataLoader(test_datasetC, batch_size=50)
```

```
In [ ]: fnnc=feedForward(3,3000)
fnnc
```

```
In [ ]: train(train_setC, test_setC, fnnc, 20, concat=True)
```

Installed Packages

```
In [1]: import sys
!{sys.executable} -m pip install contractions
!{sys.executable} -m pip install gensim==4.2.0
!pip install scikit-learn
!pip install torch torchvision torchaudio
Looking in indexes: https://pypi.org/simple, (https://pypi.org/simple,) https://us-python.pkg.dev/colab-wheels/public/simple/ (https://us-python.pkg.dev/colab-wheels/public/simple/)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.8/dist-packages (1.2.1)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.8/dist-packages (from scikit-learn) (1.10.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.8/dist-packages (from scikit-learn) (3.1.0)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.8/dist-packages (from scikit-learn) (1.2.0)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.8/dist-packages (from scikit-learn) (1.24.2)
Looking in indexes: https://pypi.org/simple, (https://pypi.org/simple,) https://us-python.pkg.dev/colab-wheels/public/simple/ (https://us-python.pkg.dev/colab-wheels/public/simple/)
Requirement already satisfied: torch in /usr/local/lib/python3.8/dist-packages (1.13.1+cu116)
Requirement already satisfied: torchvision in /usr/local/lib/python3.8/dist-packages (0.14.1+cu116)
Requirement already satisfied: torchaudio in /usr/local/lib/python3.8/dist-packages (0.13.1+cu116)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.8/dist-packages (from torch) (4.5.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.8/dist-packages (from torchvision) (1.24.2)
Requirement already satisfied: requests in /usr/local/lib/python3.8/dist-packages (from torchvision) (2.25.1)
```

```
In [2]: ## Importing and installing libraries

import numpy as np
import copy
import pandas as pd
import warnings
import re
import sys
import nltk
from gensim.models import Word2Vec
from nltk.corpus import stopwords
import string
from torch import nn
import torch.nn as nn
import torch.nn.functional as F
import torch
from torch.nn import CrossEntropyLoss, Softmax, Linear
from torch.optim import SGD, Adam
from sklearn.metrics.pairwise import cosine_similarity
from torch.optim.lr_scheduler import ReduceLROnPlateau
from nltk.stem import WordNetLemmatizer
from gensim.models import KeyedVectors
from gensim import utils
from scipy.sparse import hstack
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Perceptron
from sklearn.metrics import classification_report
from sklearn import svm
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import LinearSVC
from statistics import mean
from os import path
import os.path
import gensim
import gensim.downloader
from sklearn.svm import LinearSVC

nltk.download('punkt')

warnings.filterwarnings('ignore')

import contractions

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
```

1. Dataset Generation

```
In [3]: from google.colab import drive
drive.mount('/content/drive/')
%cd /content/drive/My Drive/Colab Notebooks/
```

Mounted at /content/drive/
/content/drive/My Drive/Colab Notebooks

```
In [4]: #fields required in the balanced dataframe from the original dataset
input_column=["review_body", "star_rating"]

#reading the original dataset to filter the columns that are required
input_df =pd.read_csv('./amazon_reviews_us_Beauty_v1_00.tsv',usecols=input_column,sep='\t',error_bad_lines=False)
```

```
In [5]: #Creating 3 different classes to get 20000 data from each class to avoid computational burden

class_one_df =(input_df[(input_df['star_rating'] == 1) | (input_df['star_rating'] == 2) ]).sample(n=20000)
class_one_df['class']=1

class_two_df =(input_df[(input_df['star_rating'] == 3)]).sample(n=20000)
class_two_df['class']=2

class_three_df =(input_df[(input_df['star_rating'] == 4) | (input_df['star_rating'] == 5) ]).sample(n=20000)
class_three_df['class']=3

#Combining all the data received from each class into a single balanced dataframe

amazon_balanced_df = pd.concat([class_one_df, class_two_df, class_three_df])

#Resetting the index as we have retrieved different data according to the classes created.
#Therefore, we will have irregular or unsorted index keys.
#We will reset the index to the new and incremental values from 0

amazon_balanced_df = amazon_balanced_df.reset_index(drop=True)

# Created a new dataframe consisting of the two columns (star_rating and review_body)
#along with class one assigned to them on the basis of star_rating. We are also resetting the index
```

Data Cleaning

Handling null values

```
In [6]: #We are changing all null values to an empty string

amazon_balanced_df = amazon_balanced_df.fillna('')
```

```
In [7]: #Uncleaned data copy
amazon_df=amazon_balanced_df.copy()
```

Convert all reviews into lowercase

```
In [8]: # Converting all review body into lowercase

amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.lower()
```

Remove the HTML from the reviews

```
In [9]: # Removing all the html tags from each review body

amazon_balanced_df['review_body']=amazon_balanced_df['review_body'].apply(lambda x : re.sub('<.*?>','',str(x)))
```

Remove the URLs from the reviews

```
In [10]: # Removing all the URLs from each review body

amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].apply(lambda y: re.split('https://\/.*', y)[0])
```

Remove non-alphabetical characters

```
In [11]: # Removing all the non alphabetic chaarcters(symbols, numbers) from each review body

amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].apply(lambda z: " ".join([re.sub('[^A-Za-
```

Remove extra spaces

```
In [12]: # Will remove leading and trailing spaces

amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.strip()
```

Perform contractions on the review_body

```
In [13]: ## This will elongate the short form used in sentences like (I'll ---> I will)

amazon_balanced_df['without_contraction'] = amazon_balanced_df['review_body'].apply(lambda a: [contractions.fix
amazon_balanced_df['review_body'] = [ ' '.join(map(str, x)) for x in amazon_balanced_df['without_contraction']]
```

Remove Punctuations

```
In [14]: amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.replace(r'[^\\w\\s]+', '')
```

2. Word Embedding

(a) Downloading pretrained word2vec-google-news-300

```
In [ ]: # word2vec_model = gensim.downloader.load('word2vec-google-news-300')
```

```
In [ ]: # word2vec_model.save('Gensim_word2vec_model.kv')
```

```
In [15]: from gensim.models import KeyedVectors
word2vec_model= KeyedVectors.load("Gensim_word2vec_model.kv")
```

(b) Training word2vec model on our own dataset

```
In [16]: class dataEmbed:
    def __init__(self, data_set):
        self.data_set = data_set

    def __iter__(self):
        for x in self.data_set:
            yield utils.simple_preprocess(x)
```

```
In [ ]: # sentence_embed = dataEmbed(amazon_balanced_df.review_body)
# # window=13
# # vector_size=300
# # min_count=9
# embed_word2vec = Word2Vec(sentences=sentence_embed, vector_size=300, min_count=9, window=13)
# model = embed_word2vec.wv
```


Process to extract word2vec embeddings

In [18]: *### To concatenate first 10 Word2Vec vectors for each review as the input feature*

```
embedding_space_concat = []
for i in range(60000):
    vectorWord = np.zeros((1,300)) # change the size of the vector
    listword = amazon_df['review_body'][i].split(" ")
    for item in listword[:10]:
        if item in word2vec_model:
            vectorWord = np.concatenate([vectorWord, np.expand_dims(word2vec_model[item], axis=0)], axis=0)

    vectorWord = vectorWord[1:]
    if len(vectorWord)<10:
        for i in range(10 - len(vectorWord)):
            vectorWord = np.concatenate([vectorWord, np.zeros((1,300))], axis=0)
    embedding_space_concat.append(vectorWord)

embedding_dataset_concat = np.array(embedding_space_concat)
embedding_dataset_concat = embedding_dataset_concat.reshape(embedding_dataset_concat.shape[0], embedding_dataset_concat.shape[1]*10)
```

In [19]: print(embedding_dataset_concat.shape)

(60000, 3000)

In [20]: P_train, P_test, Q_train, Q_test = train_test_split(embedding_dataset_concat, amazon_df['class'], test_size=0.2)

```
Q_train = Q_train.reset_index(drop=True)
Q_test = Q_test.reset_index(drop=True)
```

```
print(P_train.shape, P_test.shape, Q_train.shape, Q_test.shape)
```

(48000, 3000) (12000, 3000) (48000,) (12000,)

In []: type(P_train)

Out[25]: numpy.ndarray

4. Feedforward Neural Networks

In [21]: from torch.utils.data import Dataset, DataLoader

In [22]: *#Creating a dataloader using torch*

```
class dataloader(torch.utils.data.Dataset):
    def __init__(self, dataset_record, label_record):
        self.dataset = dataset_record
        self.labels = label_record

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, index):
        dataset = self.dataset[index]
        labels = self.labels[index]

        return dataset, labels
```

In [23]: *#Creating classes to define the architecure*

```
class feedForward(nn.Module):
    def __init__(self, output_size, input_size):
        super(feedForward, self).__init__()
        self.layer1 = nn.Linear(input_size, 300)
        self.relu1 = nn.ReLU()
        self.layer2 = nn.Linear(300, 100)
        self.relu2 = nn.ReLU()
        self.layer3 = nn.Linear(100, output_size)

    def forward(self, x):
        return self.layer3(self.relu2(self.layer2(self.relu1(self.layer1(x)))))
```



```
In [24]: fnn=feedForward(3,3000)
fnn
```

```
Out[24]: feedForward(
  (layer1): Linear(in_features=3000, out_features=300, bias=True)
  (relu1): ReLU()
  (layer2): Linear(in_features=300, out_features=100, bias=True)
  (relu2): ReLU()
  (layer3): Linear(in_features=100, out_features=3, bias=True)
)
```

(b)

```
In [25]: # Convert P_train and P_test to float32
A_word2vec_train = P_train.astype(np.float32)
A_word2vec_test  = P_test.astype(np.float32)

# Subtract 1 from B_train and B_test values
B_train = Q_train - 1
B_test  = Q_test  - 1

# Create PyTorch DataLoader objects for the training and testing sets
train_dataset = dataloader(A_word2vec_train, B_train)
train_set = torch.utils.data.DataLoader(train_dataset, batch_size=50)

test_dataset = dataloader(A_word2vec_test, B_test)
test_set = torch.utils.data.DataLoader(test_dataset, batch_size=50)
```

```
In [26]: from sklearn.metrics import accuracy_score, f1_score
```

```

In [27]: f train(reviews_dataloader_train, reviews_dataloader_test, model, num_epochs, concat=True, rnn=False, gru=False,
y_pred_label_train = []
y_true_label_train = []
y_pred_label_test = []
y_true_label_test = []

# Set the device for the model
# device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
# model.to(device)

# Define the loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = Adam(model.parameters(), lr=0.001)
softmax = Softmax(dim=1)

# Define the scheduler
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=5, gamma=0.1)

# Keep track of the best model
best_model_wts = copy.deepcopy(model.state_dict())
best_acc = 0.0

# Keep track of the previous loss
loss_min = prev_loss

# Train the model
for epoch in range(num_epochs):
    print('\n Epoch: {}'.format(epoch))

    # print(reviews_dataloader_train)
    for j, (x, y) in enumerate(reviews_dataloader_train):
        y_pred = model(x)
        y_pred_label_train.append(torch.argmax(softmax(y_pred.detach()), axis=1))
        y_true_label_train.append(y.detach())
        loss = criterion(y_pred, y)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # if j % 100 == 0:
        #     print('Epoch {:03} Batch {:03}/{:03} Loss: {:.4f}'.format(epoch, j, len(reviews_dataloader_train), loss.item()))

    # Evaluate the model on the test set
    with torch.no_grad():
        for x, y in reviews_dataloader_test:
            y_pred = model(x)
            y_pred_label_test.append(torch.argmax(softmax(y_pred.detach()), axis=1))
            y_true_label_test.append(y.detach())

    # Calculate accuracy and f1-score
    y_pred_train = torch.cat(y_pred_label_train)
    y_true_train = torch.cat(y_true_label_train)
    y_pred_test = torch.cat(y_pred_label_test)
    y_true_test = torch.cat(y_true_label_test)

    train_acc = accuracy_score(y_true_train.cpu().numpy(), y_pred_train.cpu().numpy())
    test_acc = accuracy_score(y_true_test.cpu().numpy(), y_pred_test.cpu().numpy())
    train_f1 = f1_score(y_true_train.cpu().numpy(), y_pred_train.cpu().numpy(), average='macro')
    test_f1 = f1_score(y_true_test.cpu().numpy(), y_pred_test.cpu().numpy(), average='macro')

    print('Epoch: {:03}, Loss: {:.4f}, Train Acc: {:.4f}, Test Acc: {:.4f}'.format(epoch, loss.item(), train_acc, test_acc))

    # Update the learning rate
    scheduler.step()

    # Save the best model based on test accuracy
    if test_acc > best_acc:
        best_acc = test_acc
        best_model_wts = copy.deepcopy(model.state_dict())

    # Save the model checkpoint
    # if loss.item() < loss_min:
    #     print(f'Loss decreased from {loss_min:.4f} to {loss.item():.4f}. Saving model...')
    #     torch.save(model.state_dict(), 'model_checkpoint.pt')
    #     loss

```

```
In [28]: train(train_set, test_set, fnn, 20)
```

Epoch: 0
Epoch: 000, Loss: 0.9506, Train Acc: 0.5246, Test Acc: 0.5460

Epoch: 1
Epoch: 001, Loss: 0.8462, Train Acc: 0.5612, Test Acc: 0.5457

Epoch: 2
Epoch: 002, Loss: 0.7152, Train Acc: 0.6004, Test Acc: 0.5366

Epoch: 3
Epoch: 003, Loss: 0.5106, Train Acc: 0.6424, Test Acc: 0.5287

Epoch: 4
Epoch: 004, Loss: 0.4892, Train Acc: 0.6795, Test Acc: 0.5248

Epoch: 5
Epoch: 005, Loss: 0.5019, Train Acc: 0.7056, Test Acc: 0.5249

Epoch: 6
Epoch: 006, Loss: 0.4329, Train Acc: 0.7307, Test Acc: 0.5245

Epoch: 7
Epoch: 007, Loss: 0.3555, Train Acc: 0.7528, Test Acc: 0.5241

Epoch: 8
Epoch: 008, Loss: 0.2691, Train Acc: 0.7723, Test Acc: 0.5237

Epoch: 9
Epoch: 009, Loss: 0.1962, Train Acc: 0.7894, Test Acc: 0.5233

Epoch: 10
Epoch: 010, Loss: 0.1898, Train Acc: 0.8041, Test Acc: 0.5228

Epoch: 11
Epoch: 011, Loss: 0.1857, Train Acc: 0.8167, Test Acc: 0.5224

Epoch: 12
Epoch: 012, Loss: 0.1792, Train Acc: 0.8277, Test Acc: 0.5219

Epoch: 13
Epoch: 013, Loss: 0.1708, Train Acc: 0.8371, Test Acc: 0.5214

Epoch: 14
Epoch: 014, Loss: 0.1628, Train Acc: 0.8455, Test Acc: 0.5211

Epoch: 15
Epoch: 015, Loss: 0.1601, Train Acc: 0.8529, Test Acc: 0.5207

Epoch: 16
Epoch: 016, Loss: 0.1590, Train Acc: 0.8594, Test Acc: 0.5204

Epoch: 17
Epoch: 017, Loss: 0.1580, Train Acc: 0.8653, Test Acc: 0.5202

Epoch: 18
Epoch: 018, Loss: 0.1570, Train Acc: 0.8705, Test Acc: 0.5200

Epoch: 19
Epoch: 019, Loss: 0.1561, Train Acc: 0.8752, Test Acc: 0.5198

```
In [ ]:
```

Installed Packages

```
In [1]: import sys
!{sys.executable} -m pip install contractions
!{sys.executable} -m pip install gensim==4.2.0
!pip install scikit-learn
!pip install torch torchvision torchaudio
Looking in indexes: https://pypi.org/simple, (https://pypi.org/simple,) https://us-python.pkg.dev/colab-wheels/public/simple/ (https://us-python.pkg.dev/colab-wheels/public/simple/)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.8/dist-packages (1.2.1)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.8/dist-packages (from scikit-learn) (1.22.4)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.8/dist-packages (from scikit-learn) (3.1.0)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.8/dist-packages (from scikit-learn) (1.10.1)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.8/dist-packages (from scikit-learn) (1.2.0)
Looking in indexes: https://pypi.org/simple, (https://pypi.org/simple,) https://us-python.pkg.dev/colab-wheels/public/simple/ (https://us-python.pkg.dev/colab-wheels/public/simple/)
Requirement already satisfied: torch in /usr/local/lib/python3.8/dist-packages (1.13.1+cu116)
Requirement already satisfied: torchvision in /usr/local/lib/python3.8/dist-packages (0.14.1+cu116)
Requirement already satisfied: torchaudio in /usr/local/lib/python3.8/dist-packages (0.13.1+cu116)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.8/dist-packages (from torch) (4.5.0)
Requirement already satisfied: requests in /usr/local/lib/python3.8/dist-packages (from torchvision) (2.25.1)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.8/dist-packages (from torchvision)
```

```
In [2]: ## Importing and installing libraries

import numpy as np
import copy
import pandas as pd
import warnings
import re
import sys
import nltk
from gensim.models import Word2Vec
from nltk.corpus import stopwords
import string
from torch import nn
import torch.nn as nn
import torch.nn.functional as F
import torch
from torch.nn import CrossEntropyLoss, Softmax, Linear
from torch.optim import SGD, Adam
from sklearn.metrics.pairwise import cosine_similarity
from torch.optim.lr_scheduler import ReduceLROnPlateau
from nltk.stem import WordNetLemmatizer
from gensim.models import KeyedVectors
from gensim import utils
from scipy.sparse import hstack
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Perceptron
from sklearn.metrics import classification_report
from sklearn import svm
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import LinearSVC
from statistics import mean
from os import path
import os.path
import gensim
import gensim.downloader
from sklearn.svm import LinearSVC

nltk.download('punkt')

warnings.filterwarnings('ignore')

import contractions
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
```

1. Dataset Generation

```
In [3]: from google.colab import drive
drive.mount('/content/drive/')
# %cd /content/drive/My Drive/Colab Notebooks/
```

Mounted at /content/drive/

```
In [4]: %cd /content/drive/My Drive/
```

/content/drive/My Drive

```
In [5]: #fields required in the balanced dataframe from the original dataset
input_column=["review_body","star_rating"]

#reading the original dataset to filter the columns that are required
input_df =pd.read_csv('./amazon_reviews_us_Beauty_v1_00.tsv',usecols=input_column,sep='\t',error_bad_lines=False)
```

```
In [6]: #Creating 3 different classes to get 20000 data from each class to avoid computational burden

class_one_df =(input_df[(input_df['star_rating'] == 1) | (input_df['star_rating'] == 2) ]).sample(n=20000)
class_one_df['class']=1

class_two_df =(input_df[(input_df['star_rating'] == 3)]).sample(n=20000)
class_two_df['class']=2

class_three_df =(input_df[(input_df['star_rating'] == 4) | (input_df['star_rating'] == 5) ]).sample(n=20000)
class_three_df['class']=3

#Combining all the data received from each class into a single balanced dataframe

amazon_balanced_df = pd.concat([class_one_df, class_two_df, class_three_df])

#Resetting the index as we have retrieved different data according to the classes created.
#Therefore, we will have irregular or unsorted index keys.
#We will reset the index to the new and incremental values from 0

amazon_balanced_df = amazon_balanced_df.reset_index(drop=True)

# Created a new dataframe consisting of the two columns (star_rating and review_body)
#along with class one assigned to them on the basis of star_rating. We are also resetting the index
```

Data Cleaning

Handling null values

```
In [7]: #We are changing all null values to an empty string
```

```
amazon_balanced_df = amazon_balanced_df.fillna('')
```

```
In [8]: #Uncleaned data copy
```

```
amazon_df=amazon_balanced_df.copy()
```

Convert all reviews into lowercase

```
In [9]: # Converting all review body into lowercase
```

```
amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.lower()
```

Remove the HTML from the reviews

```
In [10]: # Removing all the html tags from each review body
```

```
amazon_balanced_df['review_body']=amazon_balanced_df['review_body'].apply(lambda x : re.sub('<.*?>','',str(x)))
```

Remove the URLs from the reviews

```
In [11]: # Removing all the URLs from each review body
```

```
amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].apply(lambda y: re.split('https://\\/.*', y)[0])
```

Remove non-alphabetical characters

```
In [12]: # Removing all the non alphabetic chaarcters(symbols, numbers) from each review body

amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].apply(lambda z: " ".join([re.sub('[^A-Za-
```

Remove extra spaces

```
In [13]: # Will remove leading and trailing spaces

amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.strip()
```

Perform contractions on the review_body

```
In [14]: ## This will elongate the short form used in sentences like (I'll ---> I will)

amazon_balanced_df['without_contraction'] = amazon_balanced_df['review_body'].apply(lambda a: [contractions.fix
amazon_balanced_df['review_body'] = [' '.join(map(str, x)) for x in amazon_balanced_df['without_contraction']]
```

Remove Punctuations

```
In [15]: amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.replace(r'[\w\s]+', '')
```

2. Word Embedding

(a) Downloading pretrained word2vec-google-news-300

```
In [21]: # word2vec_model = gensim.downloader.load('word2vec-google-news-300')
# word2vec_model.save('Gensim_word2vec_model.kv')

[=====] 100.0% 1662.8/1662.8MB downloaded
```

```
In [16]: from gensim.models import KeyedVectors
word2vec_model= KeyedVectors.load("Gensim_word2vec_model.kv")
```

Process to extract word2vec embeddings

```
In [ ]: # embedding_space_concat = []
# for i in range(60000):
#     vectorWord = np.zeros((1,300)) # change the size of the vector
#     listword = amazon_df['review_body'][i].split(" ")
#     for item in listword[:20]:
#         if item in word2vec_model:
#             vectorWord = np.concatenate([vectorWord, np.expand_dims(word2vec_model[item], axis=0)], axis=0)

#     vectorWord = vectorWord[1:]
#     if len(vectorWord)<20:
#         for i in range(20 - len(vectorWord)):
#             vectorWord = np.concatenate([vectorWord, np.zeros((1,300))], axis=0)
#     embedding_space_concat.append(vectorWord)

# embedding_dataset_concat = np.array(embedding_space_concat)
# embedding_dataset_concat = embedding_dataset_concat.reshape(embedding_dataset_concat.shape[0], embedding_data
# embedding_dataset_concat
```



```
In [17]: embedding_space_concat = []
for i in range(60000):
    vectorWord = [] # change the size of the vector
    listword = amazon_df['review_body'][i].split(" ")
    for item in listword[:20]:
        if item in word2vec_model:
            x=np.reshape(word2vec_model[item], (1, 300))
            vectorWord.append(x)
    vectorWord=vectorWord[1:]
    if len(vectorWord) < 20:
        di = 20 - len(vectorWord)
        vectorWord += [np.zeros((1, 300))] * di

    embedding_space_concat.append(vectorWord)
embedding_dataset_concat=np.array(embedding_space_concat)
embedding_dataset_concat=embedding_dataset_concat.reshape(embedding_dataset_concat.shape[0], embedding_dataset_concat.shape[1], embedding_dataset_concat.shape[2])
```

```
In [18]: embedding_dataset_concat.shape
```

```
Out[18]: (60000, 20, 300)
```

```
In [19]: A_train, A_test, B_train, B_test = train_test_split(embedding_dataset_concat, amazon_df['class'], test_size=0.2)

B_train = B_train.reset_index(drop=True)
B_test = B_test.reset_index(drop=True)
```

5. Recurrent Neural Networks

```
In [21]: from torch.utils.data import Dataset, DataLoader
```

```
In [22]: #Creating a dataloader using torch
class dataloader(torch.utils.data.Dataset):
    def __init__(self, dataset_record, label_record):
        self.dataset = dataset_record
        self.labels = label_record

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, index):
        dataset = self.dataset[index]
        labels = self.labels[index]

        return dataset, labels
```

```
In [23]: class RNN(nn.Module):
    def __init__(self, classes, layer, batch_size):
        super(RNN, self).__init__()
        self.rnn = nn.RNN(300, 600, layer, batch_first=True)
        self.h1 = torch.randn(layer, batch_size, 600)
        self.linear = nn.Linear(600, classes)

    def forward(self, x):
        return self.linear(self.rnn(x)[0][:, -1])
```

```
In [24]: rnn_m = RNN(3, 2, 100)
rnn_m
```

```
Out[24]: RNN(
  (rnn): RNN(300, 600, num_layers=2, batch_first=True)
  (linear): Linear(in_features=600, out_features=3, bias=True)
)
```

5. (a) Non Gated RNN

```
In [25]: # Convert A_train and A_test to float32
A_word2vec_train = A_train.astype(np.float32)
A_word2vec_test = A_test.astype(np.float32)

# Subtract 1 from B_train and B_test values
B_train = B_train - 1
B_test = B_test - 1

# Create PyTorch DataLoader objects for the training and testing sets
train_dataset = dataloader(A_word2vec_train, B_train)
train_set = torch.utils.data.DataLoader(train_dataset, batch_size=100)

test_dataset = dataloader(A_word2vec_test, B_test)
test_set = torch.utils.data.DataLoader(test_dataset, batch_size=100)
```



```
In [26]: from sklearn.metrics import accuracy_score, f1_score
```

```
In [27]: def train(reviews_dataloader_train, reviews_dataloader_test, model, num_epochs, concat=True, rnn=True, gru=False):
    y_pred_label_train = []
    y_true_label_train = []
    y_pred_label_test = []
    y_true_label_test = []

    # Set the device for the model
    # device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    # model.to(device)

    # Define the loss function and optimizer
    criterion = nn.CrossEntropyLoss()
    optimizer = Adam(model.parameters(), lr=0.001)
    # optimizer = SGD(rnn.parameters(), lr=1e-2)
    scheduler = ReduceLROnPlateau(optimizer)

    # optimizer = Adam(model.parameters(), lr=0.001)
    softmax = Softmax(dim=1)

    # Define the scheduler
    scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=5, gamma=0.1)

    # Keep track of the best model
    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0

    # Keep track of the previous loss
    loss_min = prev_loss

    # Train the model
    for epoch in range(num_epochs):
        print('\n Epoch: {}'.format(epoch))

        # print(reviews_dataloader_train)
        for j, (x, y) in enumerate(reviews_dataloader_train):
            y_pred = model(x)
            y_pred_label_train.append(torch.argmax(softmax(y_pred.detach()), axis=1))
            y_true_label_train.append(y.detach())
            loss = criterion(y_pred, y)

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            # if j % 100 == 0:
            #     print('Epoch {:03} Batch {:03}/{:03} Loss: {:.4f}'.format(epoch, j, len(reviews_dataloader_train), loss.item()))

        # Evaluate the model on the test set
        with torch.no_grad():
            for x, y in reviews_dataloader_test:
                y_pred = model(x)
                y_pred_label_test.append(torch.argmax(softmax(y_pred.detach()), axis=1))
                y_true_label_test.append(y.detach())

        # Calculate accuracy and f1-score
        y_pred_train = torch.cat(y_pred_label_train)
        y_true_train = torch.cat(y_true_label_train)
        y_pred_test = torch.cat(y_pred_label_test)
        y_true_test = torch.cat(y_true_label_test)

        train_acc = accuracy_score(y_true_train.cpu().numpy(), y_pred_train.cpu().numpy())
        test_acc = accuracy_score(y_true_test.cpu().numpy(), y_pred_test.cpu().numpy())
        train_f1 = f1_score(y_true_train.cpu().numpy(), y_pred_train.cpu().numpy(), average='macro')
        test_f1 = f1_score(y_true_test.cpu().numpy(), y_pred_test.cpu().numpy(), average='macro')

        print('Epoch: {:03}, Loss: {:.4f}, Train Acc: {:.4f}, Test Acc: {:.4f}'.format(epoch, loss.item(), train_acc, test_acc))

        # Update the learning rate
        scheduler.step()

        # Save the best model based on test accuracy
        if test_acc > best_acc:
            best_acc = test_acc
            best_model_wts = copy.deepcopy(model.state_dict())

        # Save the model checkpoint
        # if loss.item() < loss_min:
        #     print(f'Loss decreased from {loss_min:.4f} to {loss.item():.4f}. Saving model...')
        #     torch.save(model.state_dict(), 'model_checkpoint.pt')
        #     loss
```

```
In [28]: train(train_set, test_set, rnn_m, 20)
```

```
Epoch: 0
Epoch: 000, Loss: 1.0730, Train Acc: 0.3482, Test Acc: 0.3767

Epoch: 1
Epoch: 001, Loss: 1.1250, Train Acc: 0.3590, Test Acc: 0.3447

Epoch: 2
Epoch: 002, Loss: 1.1038, Train Acc: 0.3572, Test Acc: 0.3496

Epoch: 3
Epoch: 003, Loss: 1.1070, Train Acc: 0.3566, Test Acc: 0.3521

Epoch: 4
Epoch: 004, Loss: 1.1113, Train Acc: 0.3562, Test Acc: 0.3537

Epoch: 5
Epoch: 005, Loss: 1.0876, Train Acc: 0.3581, Test Acc: 0.3576

Epoch: 6
Epoch: 006, Loss: 1.0864, Train Acc: 0.3598, Test Acc: 0.3606

Epoch: 7
Epoch: 007, Loss: 1.0828, Train Acc: 0.3612, Test Acc: 0.3629

Epoch: 8
Epoch: 008, Loss: 1.0826, Train Acc: 0.3623, Test Acc: 0.3647

Epoch: 9
Epoch: 009, Loss: 1.0824, Train Acc: 0.3632, Test Acc: 0.3661

Epoch: 10
Epoch: 010, Loss: 1.0850, Train Acc: 0.3645, Test Acc: 0.3672

Epoch: 11
Epoch: 011, Loss: 1.0849, Train Acc: 0.3657, Test Acc: 0.3682

Epoch: 12
Epoch: 012, Loss: 1.0849, Train Acc: 0.3666, Test Acc: 0.3690

Epoch: 13
Epoch: 013, Loss: 1.0849, Train Acc: 0.3674, Test Acc: 0.3697

Epoch: 14
Epoch: 014, Loss: 1.0849, Train Acc: 0.3682, Test Acc: 0.3703

Epoch: 15
Epoch: 015, Loss: 1.0849, Train Acc: 0.3689, Test Acc: 0.3709

Epoch: 16
Epoch: 016, Loss: 1.0850, Train Acc: 0.3695, Test Acc: 0.3713

Epoch: 17
Epoch: 017, Loss: 1.0850, Train Acc: 0.3701, Test Acc: 0.3717

Epoch: 18
Epoch: 018, Loss: 1.0850, Train Acc: 0.3706, Test Acc: 0.3721

Epoch: 19
Epoch: 019, Loss: 1.0850, Train Acc: 0.3710, Test Acc: 0.3725
```

5. (b) Gated RNN

```
In [29]: class GatedRNN(nn.Module):
          def __init__(self, num_classes, layers, batch_size):
              super(GatedRNN, self).__init__()
              self.rnn = nn.GRU(300, 300, layers, batch_first=True)
              self.h1 = torch.randn(layers, batch_size, 300)
              self.linear = nn.Linear(300, num_classes)

          def forward(self, x):
              return self.linear(self.rnn(x)[0][:, -1])
```

```
In [30]: gru = GatedRNN(5, 1, 100)
```

```
In [31]: train(train_set, test_set, gru, 20, True, True, True)
```

```
Epoch: 0
Epoch: 000, Loss: 0.9270, Train Acc: 0.5037, Test Acc: 0.5558

Epoch: 1
Epoch: 001, Loss: 0.8748, Train Acc: 0.5415, Test Acc: 0.5694

Epoch: 2
Epoch: 002, Loss: 0.8604, Train Acc: 0.5606, Test Acc: 0.5769

Epoch: 3
Epoch: 003, Loss: 0.8465, Train Acc: 0.5739, Test Acc: 0.5822

Epoch: 4
Epoch: 004, Loss: 0.8341, Train Acc: 0.5850, Test Acc: 0.5863

Epoch: 5
Epoch: 005, Loss: 0.8319, Train Acc: 0.5962, Test Acc: 0.5909

Epoch: 6
Epoch: 006, Loss: 0.8230, Train Acc: 0.6048, Test Acc: 0.5939

Epoch: 7
Epoch: 007, Loss: 0.8148, Train Acc: 0.6116, Test Acc: 0.5961

Epoch: 8
Epoch: 008, Loss: 0.8064, Train Acc: 0.6173, Test Acc: 0.5978

Epoch: 9
Epoch: 009, Loss: 0.7969, Train Acc: 0.6221, Test Acc: 0.5992

Epoch: 10
Epoch: 010, Loss: 0.7995, Train Acc: 0.6266, Test Acc: 0.6007

Epoch: 11
Epoch: 011, Loss: 0.7995, Train Acc: 0.6304, Test Acc: 0.6019

Epoch: 12
Epoch: 012, Loss: 0.7987, Train Acc: 0.6336, Test Acc: 0.6029

Epoch: 13
Epoch: 013, Loss: 0.7977, Train Acc: 0.6364, Test Acc: 0.6037

Epoch: 14
Epoch: 014, Loss: 0.7966, Train Acc: 0.6389, Test Acc: 0.6044

Epoch: 15
Epoch: 015, Loss: 0.7967, Train Acc: 0.6411, Test Acc: 0.6050

Epoch: 16
Epoch: 016, Loss: 0.7967, Train Acc: 0.6430, Test Acc: 0.6056

Epoch: 17
Epoch: 017, Loss: 0.7967, Train Acc: 0.6447, Test Acc: 0.6061

Epoch: 18
Epoch: 018, Loss: 0.7967, Train Acc: 0.6463, Test Acc: 0.6065

Epoch: 19
Epoch: 019, Loss: 0.7967, Train Acc: 0.6476, Test Acc: 0.6069
```

Installed Packages

```
In [ ]: ▶ import sys
        !{sys.executable} -m pip install contractions
        !{sys.executable} -m pip install gensim==4.2.0
        !pip install scikit-learn
        !pip install torch torchvision torchaudio
```

```
In [ ]: ## Importing and installing Libraries

import numpy as np
import copy
import pandas as pd
import warnings
import re
import sys
import nltk
from gensim.models import Word2Vec
from nltk.corpus import stopwords
import string
from torch import nn
import torch.nn as nn
import torch.nn.functional as F
import torch
from torch.nn import CrossEntropyLoss, Softmax, Linear
from torch.optim import SGD, Adam
from sklearn.metrics.pairwise import cosine_similarity
from torch.optim.lr_scheduler import ReduceLROnPlateau
from nltk.stem import WordNetLemmatizer
from gensim.models import KeyedVectors
from gensim import utils
from scipy.sparse import hstack
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Perceptron
from sklearn.metrics import classification_report
from sklearn import svm
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import LinearSVC
from statistics import mean
from os import path
import os.path
import gensim
import gensim.downloader
from sklearn.svm import LinearSVC

nltk.download('punkt')

warnings.filterwarnings('ignore')

import contractions
```

1. Dataset Generation

In [6]: **▶** *#fields required in the balanced dataframe from the original dataset*

#reading the original dataset to filter the columns that are required
input_df = pd.read_csv('https://s3.amazonaws.com/amazon-reviews-pds/tsv/ama

In [8]: **▶** *#Creating 3 different classes to get 20000 data from each class to avoid c*

class_one_df =(input_df[(input_df['star_rating'] == 1) | (input_df['star_r
class_one_df['class']=1

class_two_df =(input_df[(input_df['star_rating'] == 3)]).sample(n=20000)
class_two_df['class']=2

class_three_df =(input_df[(input_df['star_rating'] == 4) | (input_df['star
class_three_df['class']=3

#Combining all the data received from each class into a single balanced d

amazon_balanced_df = pd.concat([class_one_df, class_two_df, class_three_df

#Resetting the index as we have retrieved different data according to the
#Therefore, we will have irregular or unsorted index keys.
#We will reset the index to the new and incremental values from 0

amazon_balanced_df = amazon_balanced_df.reset_index(drop=True)

Created a new dataframe consisting of the two columns (star_rating and r
#along with class one assigned to them on the basis of star_rating. We are

Data Cleaning

Handling null values

In [9]: **▶** *#We are changing all null values to an empty string*

amazon_balanced_df = amazon_balanced_df.fillna('')

In [10]: **▶** *#Uncleaned data copy*

amazon_df=amazon_balanced_df.copy()

Convert all reviews into lowercase

```
In [11]: ▶ # Converting all review body into lowercase  
  
amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.lower
```

Remove the HTML from the reviews

```
In [12]: ▶ # Removing all the html tags from each review body  
  
amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].apply(lambda x: re.sub('<.*>', '', x))
```

Remove the URLs from the reviews

```
In [13]: ▶ # Removing all the URLs from each review body  
  
amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].apply(lambda x: re.sub('https?://.*', '', x))
```

Remove non-alphabetical characters

```
In [14]: ▶ # Removing all the non alphabetic characters(symbols, numbers) from each review body  
  
amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].apply(lambda x: re.sub('[^a-zA-Z]', '', x))
```

Remove extra spaces

```
In [15]: ▶ # Will remove leading and trailing spaces  
  
amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.strip
```


Perform contractions on the review_body

```
In [16]: ▶ ## This will elongate the short form used in sentences like (I'll ---> I w  
amazon_balanced_df['without_contraction'] = amazon_balanced_df['review_bod  
amazon_balanced_df['review_body'] = [' '.join(map(str, x)) for x in amazon
```

Remove Punctuations

```
In [17]: ▶ amazon_balanced_df['review_body'] = amazon_balanced_df['review_body'].str.
```

2. Word Embedding

(a) Downloading pretrained word2vec-google-news-300

```
In [20]: ▶ # word2vec_model = gensim.downloader.Load('word2vec-google-news-300')  
# word2vec_model.save('Gensim_word2vec_model.kv')
```

```
In [21]: ▶ from gensim.models import KeyedVectors  
word2vec_model= KeyedVectors.load("Gensim_word2vec_model.kv")
```

Process to extract word2vec embeddings

```
In [23]: embedding_space_concat = []
for i in range(60000):
    vectorWord = [] # change the size of the vector
    listword = amazon_df['review_body'][i].split(" ")
    for item in listword[:20]:
        if item in word2vec_model:
            x=np.reshape(word2vec_model[item], (1, 300))
            vectorWord.append(x)
    vectorWord=vectorWord[1:]
    if len(vectorWord) < 20:
        di = 20 - len(vectorWord)
        vectorWord += [np.zeros((1, 300))] * di

    embedding_space_concat.append(vectorWord)
embedding_dataset_concat=np.array(embedding_space_concat)
embedding_dataset_concat=embedding_dataset_concat.reshape(embedding_dataset_concat.shape[0], embedding_dataset_concat.shape[1], embedding_dataset_concat.shape[2])
```

```
In [24]: embedding_dataset_concat.shape
```

```
Out[24]: (60000, 20, 300)
```

```
In [25]: A_train, A_test, B_train, B_test = train_test_split(embedding_dataset_concat, y_train, y_test, test_size=0.2, random_state=42)
```

```
In [26]: B_train = B_train.reset_index(drop=True)
B_test = B_test.reset_index(drop=True)

print(A_train.shape, A_test.shape, B_train.shape, B_test.shape)
```

```
(48000, 20, 300) (12000, 20, 300) (48000,) (12000,)
```

5. Recurrent Neural Networks

```
In [27]: from torch.utils.data import Dataset, DataLoader
```

```
In [28]: ▶ #Creating a DataLoader using torch
class dataloader(torch.utils.data.Dataset):
    def __init__(self, dataset_record, label_record):
        self.dataset = dataset_record
        self.labels = label_record

    def __len__(self):
        return len(self.labels)

    def __getitem__(self, index):
        dataset = self.dataset[index]
        labels = self.labels[index]

        return dataset, labels
```

```
In [29]: ▶ # Convert A_train and A_test to float32
A_word2vec_train = A_train.astype(np.float32)
A_word2vec_test = A_test.astype(np.float32)

# Subtract 1 from B_train and B_test values
B_train = B_train - 1
B_test = B_test - 1

# Create PyTorch DataLoader objects for the training and testing sets
train_dataset = dataloader(A_word2vec_train, B_train)
train_set = torch.utils.data.DataLoader(train_dataset, batch_size=100)

test_dataset = dataloader(A_word2vec_test, B_test)
test_set = torch.utils.data.DataLoader(test_dataset, batch_size=100)
```

```
In [30]: ▶ from sklearn.metrics import accuracy_score, f1_score
```



```

In [31]: ▶ def train(reviews_dataloader_train, reviews_dataloader_test, model, num_epochs):
    y_pred_label_train = []
    y_true_label_train = []
    y_pred_label_test = []
    y_true_label_test = []

    # Set the device for the model
    # device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    # model.to(device)

    # Define the loss function and optimizer
    criterion = nn.CrossEntropyLoss()
    optimizer = Adam(model.parameters(), lr=0.001)
    # optimizer = SGD(rnn.parameters(), lr=1e-2)
    scheduler = ReduceLROnPlateau(optimizer)

    # optimizer = Adam(model.parameters(), lr=0.001)
    softmax = Softmax(dim=1)

    # Define the scheduler
    scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=5, gamma=0.5)

    # Keep track of the best model
    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0

    # Keep track of the previous loss
    loss_min = prev_loss

    # Train the model
    for epoch in range(num_epochs):
        print('\n Epoch: {}'.format(epoch))

        # print(reviews_dataloader_train)
        for j, (x, y) in enumerate(reviews_dataloader_train):
            y_pred = model(x)
            y_pred_label_train.append(torch.argmax(softmax(y_pred.detach()), dim=1))
            y_true_label_train.append(y.detach())
            loss = criterion(y_pred, y)

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            # if j % 100 == 0:
            #     print('Epoch {}:03 Batch {}:03/ {}:03 Loss: {:.4f}'.format(epoch, j, num_epochs, loss))

        # Evaluate the model on the test set
        with torch.no_grad():
            for x, y in reviews_dataloader_test:
                y_pred = model(x)
                y_pred_label_test.append(torch.argmax(softmax(y_pred.detach()), dim=1))
                y_true_label_test.append(y.detach())

    # Calculate accuracy and f1-score
    y_pred_train = torch.cat(y_pred_label_train)
    y_true_train = torch.cat(y_true_label_train)

```

```

y_pred_test = torch.cat(y_pred_label_test)
y_true_test = torch.cat(y_true_label_test)

train_acc = accuracy_score(y_true_train.cpu().numpy(), y_pred_train.cpu().numpy())
test_acc = accuracy_score(y_true_test.cpu().numpy(), y_pred_test.cpu().numpy())
train_f1 = f1_score(y_true_train.cpu().numpy(), y_pred_train.cpu().numpy())
test_f1 = f1_score(y_true_test.cpu().numpy(), y_pred_test.cpu().numpy())

print('Epoch: {:03}, Loss: {:.4f}, Train Acc: {:.4f}, Test Acc: {:.4f}, Train F1: {:.4f}, Test F1: {:.4f}'.format(
    epoch, loss, train_acc, test_acc, train_f1, test_f1))

# Update the Learning rate
scheduler.step()

# Save the best model based on test accuracy
if test_acc > best_acc:
    best_acc = test_acc
    best_model_wts = copy.deepcopy(model.state_dict())

# Save the model checkpoint
# if loss.item() < loss_min:
#     print(f'Loss decreased from {loss_min:.4f} to {loss.item():.4f}')
#     torch.save(model.state_dict(), 'model_checkpoint.pt')
#     loss_min = loss.item()

```

5. (c)

```

In [32]: ▶ class LSTM(nn.Module):
    def __init__(self, num_classes, layers, hidden_size):
        super(LSTM, self).__init__()
        self.lstm = nn.LSTM(300, hidden_size, layers, batch_first=True)
        self.linear = nn.Linear(hidden_size, num_classes)

    def forward(self, x):
        return self.linear(self.lstm(x)[0][:, -1])

```

```

In [37]: ▶ lstm = LSTM(3,30,100)

```

```
In [50]: tr = train(train_set, test_set, lstm, 5, True, True)
```

```
Epoch: 0  
Epoch: 000, Loss: 0.8173, Train Acc: 0.6202, Test Acc: 0.6027  
  
Epoch: 1  
Epoch: 001, Loss: 0.7975, Train Acc: 0.6246, Test Acc: 0.6059  
  
Epoch: 2  
Epoch: 002, Loss: 0.7789, Train Acc: 0.6294, Test Acc: 0.6079  
  
Epoch: 3  
Epoch: 003, Loss: 0.7677, Train Acc: 0.6345, Test Acc: 0.6096  
  
Epoch: 4  
Epoch: 004, Loss: 0.7480, Train Acc: 0.6393, Test Acc: 0.6114  
  
Epoch: 5
```

```
In [1]: print('Accuracy for LSTM is :61.14' )
```

```
Accuracy for LSTM is :61.14
```

```
In [ ]:
```